


FastAPI and Tile38: Ultra fast retrieval of geospatial data



Ganesh N. Sivalingam



Overview

- Introduction to myself and Parkopedia
- Example problem
- Tile38
- FastAPI
- Demo!
- Benchmarking against PostGIS
- Conclusions

About me

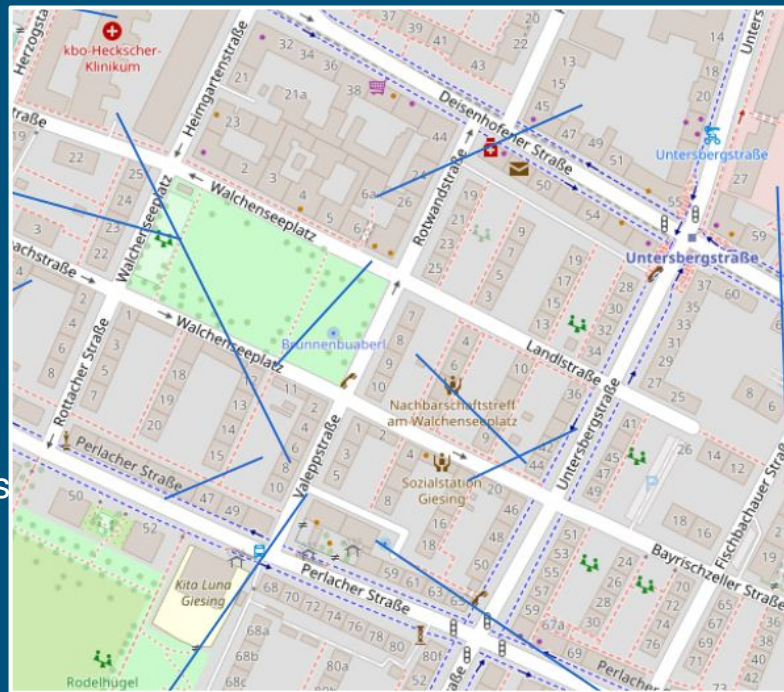
- Started at Parkopedia 6 years ago as a Data Scientist, Big Data
- Head of Data Science, technical work predominantly focussed on Data Engineering



- World's largest parking search engine, providing parking information from Apple to Lamborghini
 - Products
 - Database of static and dynamic parking spaces
 - Predictions of parking availability
 - Parking search routes
 - Many others
 - Moving into EV data
- 
- We are*
HIRING
- Machine learning modelling
 - Routing and route optimisation

Example problem

- Millions of parking locations to store
- Retrieve parking locations that are
 - Within a certain distance of a car
 - Are free or paid
 - We need to be able to serve results as quickly as possible to service the large fleet of users
- Code is available at github.com/gnsiva/tile38-fastapi




Generated dummy parking locations


Tile38

- Open source
- Fast, in-memory, geospatial database and geofencing server
- Rapid release cycle, and very quick to fix bugs
- Thoroughly documented
- Supports several geospatial queries


Nearby




Intersects



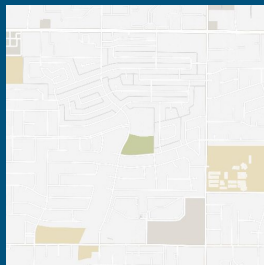
Within



Geofence



Roaming



Tile38 is easy to use

- Compatible with most Redis clients
- Docker container and binaries available so very easy to get started
- Dedicated Python package `py1e38`, but it isn't very mature so we will use `redis-py`

Data is held in a structure similar to a Python dictionary within dictionary

Inserting data

```
SET <key> <field> OBJECT <geojson>
```

Retrieving data directly

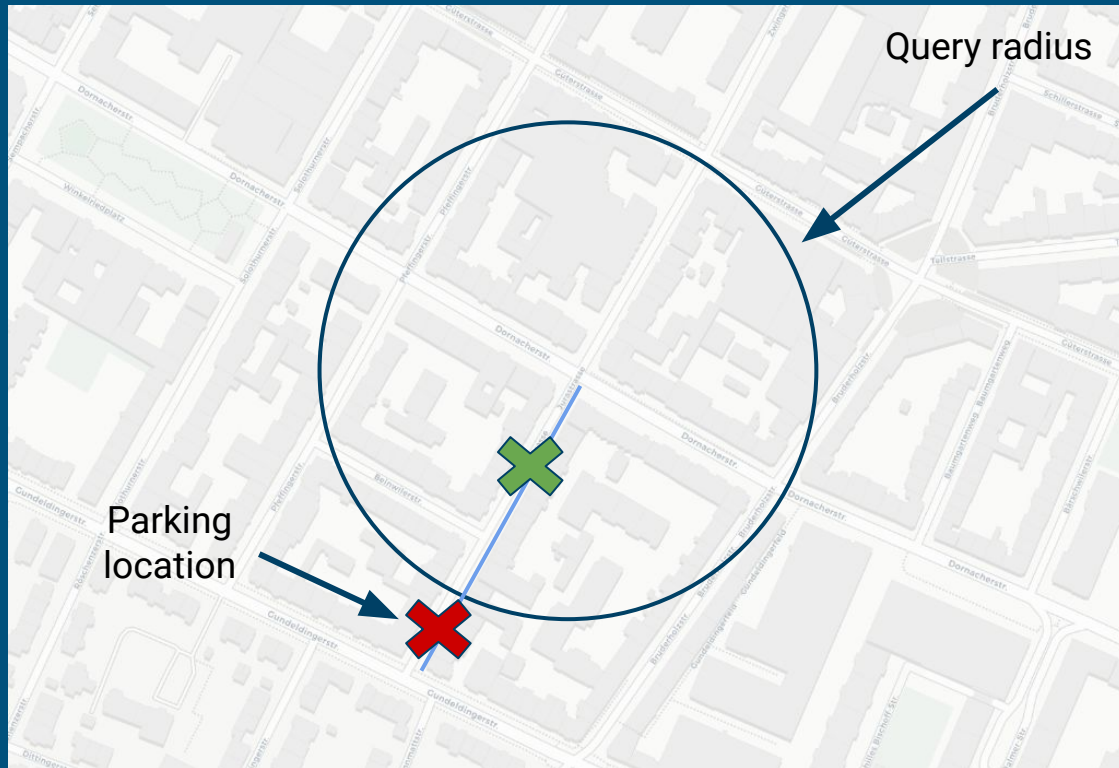
```
GET <key> <field>
```

Retrieving all objects near a location

```
NEARBY <key> POINT <lat> <lng> <meters>
```

Why not use Redis?

- Extremely fast BUT
- Can only store points, not polylines or polygons
- Limited querying
 - Radius
 - Bounding box





- An open-source “Fast”, “API” framework for Python
- Developer friendly and easy to work with
- Automatically generates OpenAPI schemas and launches documentation
- Built around Python’s asyncio module, which allows it to handle high volumes of I/O bound queries.
- Great for developing microservices

FastAPI basic app

Instantiate the FastAPI app

Decorator to
add
endpoints

Create client to
connect to
Tile38

Define query
parameters and
type validation

Run query to
retrieve parking
locations

Start the server that will
run the FastAPI app

```
app = FastAPI()


@app.get("/nearby-parking-locations")
async def get_nearby(
    latitude: float,
    longitude: float,
    d_meters: int,
):
    c = redis.Redis(host="localhost", port=9851)

    return c.execute_command(
        "NEARBY", "parking-locations", "POINT", latitude, longitude, d_meters)

if __name__ == "__main__":
    uvicorn.run("api_v1:app")
```

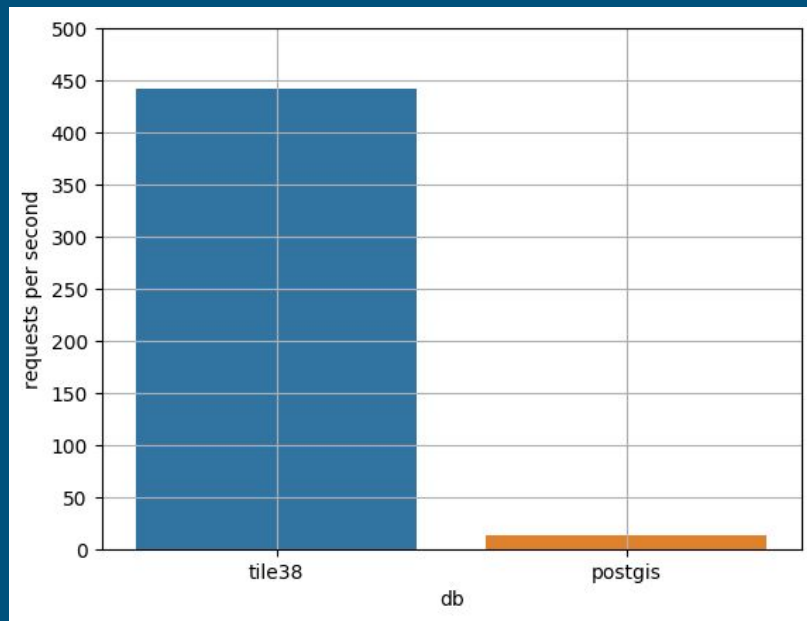


Demo



Benchmark

- Tested against PostGIS
- Spatial index on geometry column in PostGIS
- Used Apache Bench to conduct benchmarks



>30x increase in requests per second

Conclusion

- We developed a “Fast” API for retrieving parking locations
- Parkopedia
- Where to learn more
 - tile38.com
 - fastapi.tiangolo.com

Demo code and slides

github.com/gnsiva/tile38-fastapi

We are hiring Data Engineers and
Computer Vision Research Engineers

parkopedia.com/jobs

Demo code and slides

github.com/gnsiva/tile38-fastapi

We are hiring Data Engineers and
Computer Vision Research Engineers

parkopedia.com/jobs

Thank you!



Questions?